

Build Plan

The AI Marketing Engineer · Technical build plan v1 · September 24, 2026 · Sebastian van der Leegte

This document turns the product plan into an engineering sequence: the architecture, the stack, the data model, and twelve components built one at a time, each with a definition of done. It is the plan we execute component by component.

What we are building. A system a founder connects to their product, their ICP and their accounts; it decides what to build first, writes in the founder's voice, maps where the ICP pays attention, generates and indexes pages from product data, runs launches, builds relationships and tools, and reports one number every Friday. Every item on the source post is a component below, except the one that stays human.

Contents

1. Principles
2. Architecture
3. Stack
4. Data model
5. Build sequence (Components 0–12)
6. Component specifications
7. Quality gates and metrics
8. Environment, security and cost
9. First two weeks, day by day


Build order in one line: **schema** → **intake** → **scorecard** → **voice** → **attention** → **report** → **judgment** → **pages** → **launches** → **relationships** → **site/onboarding** → **tools** → **learning**.

1. Principles

1. **Outcome schema before features.** Nothing ships until every company, plan, post and weekly number lands in one schema. The dataset is the moat; it exists from company #1.
2. **Wizard of Oz is allowed, black boxes are not.** Any step may be done by a human behind the product for the first 90 days, but it must be done *through* the product so the input, output and approval are logged.
3. **The founder approves, the product ships.** No post, page or outreach goes out without a one-tap approval in the feed. Approval rate is a first-class metric.
4. **One number.** Impressions inside ICP is the north star. Every component must move it or explain why it will.
5. **Own the durable assets.** Email list, pages, and the signup-source field belong to the founder and live outside any social platform.
6. **Small, boring stack.** Postgres, a queue, one web app, one model provider behind an interface. Nothing exotic until the wedge is proven.

2. Architecture

Four layers on one Postgres database, with a job queue between the judgment layer and the execution workers. Everything the product does is a *Job* with an input, an output, an approval state and a result.



Request flow, one week of one company

1. Monday: the **sequencer** reads last week's metrics and the plan, emits this week's Jobs (10 posts, 1 reply list, 1 page batch, 0–1 launch tasks).
2. Voice engine consumes the weekly interview transcript, drafts posts as Jobs in pending.
3. Founder opens the approval feed; each tap moves a Job to approved; edits are stored as diffs (training signal).
4. Workers publish approved Jobs via integrations on schedule; Job becomes executed with the platform id.
5. Friday: **metrics pull** reads analytics per executed Job, matches engaged accounts to the ICP model, writes metric rows; report renders; share card is generated.
6. Outcome rows (plan → result) are appended; the sequencer reads them next Monday.

3. Stack

Layer	Choice	Why
Web app	Next.js 15, TypeScript, Tailwind	Fast to build the approval feed and report; SSR for public share cards and hosted pages
API and workers	Python 3.12, FastAPI, RQ on Redis	Python for the LLM and data work; RQ is simple, inspectable, enough for thousands of jobs a day
Database	Postgres 16 with pgvector	One database for everything; vectors for voice samples and ICP account similarity
Model access	One LLM interface; Claude as default, swappable	Voice quality is the product; keep the provider behind an interface so we can A/B and pin versions per founder
Enrichment	One enrichment API (Apollo/Clay-style) behind an interface	Attention map and ICP matching need company + person data; vendor is replaceable
Social	LinkedIn Marketing API, X API v2; founder-authorised OAuth only	Never operate accounts the founder doesn't control; official APIs only
Search	Google Search Console API; Bing Webmaster	Index and impression data for the page factory
Hosting for pages and tools	Subdomain per company on our infra (Vercel) or push to founder's CMS via adapter	Founder owns the domain; we host by default, export always
Auth	Clerk (or Auth.js)	Founder logs in plus OAuth token storage for integrations

Infra	Vercel (web), Fly.io or Railway (API + workers), Neon or Supabase (Postgres), Upstash (Redis)	Zero-ops for a two-person team; all have free or cheap tiers
Observability	Sentry, structured logs, one Metabase over Postgres	The internal dashboard is Metabase over the outcome tables
Repo	Monorepo: /web /api /workers /packages/schema	One schema package generates both Pydantic and TypeScript types

4. Data model

Eleven tables. The first seven exist before any feature; the rest arrive with their component. Every row carries `company_id`, `created_at`, `updated_at`.

Table	Key columns	Purpose
company	id, name, domain, stage, product_summary, product_data_sources (json), founders (json), plan_tier, status	One row per customer
founder	id, company_id, name, linkedin_handle, x_handle, oauth_tokens (encrypted), voice_profile_id	The person whose voice and accounts we use
icp	id, company_id, description, best_customers (json: 3+), firmographics (json), personas (json), embedding	Who the buyer is; refreshed from signup sources
voice_profile	id, founder_id, samples (json), rules (json: tone, banned phrases, formats allowed), embedding, version	Learned representation of how this founder writes
plan	id, company_id, week_start, scorecard (json: 5 channels × score + fix), sequence (json), targets (json), generated_by (rules llm human)	The judgment layer's output for a week
job	id, company_id, plan_id, type (post reply page launch_task outreach site_change tool), channel, input (json), output (json), state (pending approved edited rejected executed failed), scheduled_for, executed_at, platform_ref, approval_diff (json)	Every action the product proposes or takes
metric	id, company_id, job_id (nullable), date, name (impressions impressions_icp followers_icp clicks signups pages_indexed ...), value, source	Daily numbers, per job and per company

account	id, company_id, platform, handle, name, headline, company, icp_match_score, follows_count, interaction_json	The attention map: people and accounts around the ICP
page	id, company_id, template_id, slug, title, data (json), html_ref, status (draft published indexed), indexed_at	Page factory output
outcome	id, company_id, plan_id, week_start, targets (json), actuals (json), delta (json), notes	Plan vs result — the training set for judgment
signup_source	id, company_id, signup_id, answer_text, classified_channel, created_at	The "how did you hear about us" field, owned by us

```
-- Core invariant: every executed action has a measurable trail
job.state = 'executed'  => job.platform_ref IS NOT NULL
                        => metric rows for job_id appear within 7 days
plan.week_start         => exactly one outcome row per (company_id, week_start) by the
                        following Monday
```

5. Build sequence

Thirteen components, numbered from zero. Each has a week window, a dependency, and a definition of done. The order is the order of dependency and proof, not of ambition: the wedge (0–6) is what ten design partners run on; 7–12 complete the formula.

#	Component	Weeks	Depends on	Definition of done
0	Schema and skeleton	1	—	Monorepo up, Postgres migrated with the 7 core tables, API auth, job queue running a no-op job end to end, Metabase over the DB
1	Intake and ICP model	1–2	0	Founder can connect accounts (OAuth), paste 3 customers and product URLs; product summary and ICP row generated and editable; ICP embedding stored
2	Scorecard (Judgment v0)	2	1	Five-channel scorecard with a score and one fix per channel, generated from intake within 5 minutes; rendered in the app; stored as a plan row
3	Voice engine and approval feed	3–5	1	Weekly interview upload → transcript → 8–12 drafts across 12 formats → approval feed with tap-approve/edit/reject → scheduled publish via official APIs; approval diffs stored; voice-match check gating each draft

4	Attention map	4–6	1	300+ accounts per company scored for ICP fit; daily reply/engage list of 10; reply drafts in the founder's voice through the same feed
5	Metrics and Friday report	5–7	3, 4	Daily metric pull from LinkedIn, X, Search Console; impressions-inside-ICP computed from engaged accounts × ICP match; Friday email + in-app report + public share card; signup_source widget embeddable on the founder's site
6	Sequencer (Judgment v1)	7–9	2, 5	Monday re-plan from last week's outcome row; emits the week's Jobs; explicit rule set versioned in the repo; human override logged
7	Page factory	8–12	1, 5	Three page templates (product-for-segment, competitor-alternative, use-case-with-data); pages generated from product data, hosted on a company subdomain or pushed to the founder's CMS; indexing tracked
8	Launch kit	13–18	3, 6	Six-week launch calendar generated as Jobs; Product Hunt and feature-drop checklists; asset generation (copy, images brief); launch metrics rolled into the report
9	Relationship engine	19–26 (v1), 36–42 (v2)	4, 8	v1: peer companies on the product matched for joint launches, brokered inside the app. v2: influencer and peer-founder map with outreach drafts in the founder's voice, via the feed
10	Site and onboarding engine	27–34	1, 5	40-point audit automated; landing-page and signup-flow rewrites generated and shipped via CMS adapter or hosted variant; activation email sequence; signup-to-activation tracked
11	Tool factory	32–38	7, 6	Judgment proposes free tools from product data; product generates a single-page tool (calculator, scorecard, generator) with a lead-capture form; hosted and distributed as Jobs
12	Learned judgment (v2)	40–52	6, 100 companies of outcome rows	Sequencer choices proposed by a model trained/prompted on cross-company outcomes; blind A/B against the rules engine; ships when it wins on impressions-inside-ICP over 4 weeks

Wedge = components 0–6, shipped by week 9, with 7 following by week 12. That is what the ten design partners run on and what the 90-day gate (approval > 70%, one 3× result) is measured against.

6. Component specifications

Component 0 — Schema and skeleton week 1

Deliverables	Monorepo; packages/schema with the 11 tables as a single source (SQL migrations + Pydantic + TS types generated); FastAPI with auth and /companies, /jobs, /metrics; RQ worker with a heartbeat job; Next.js shell with login; Metabase connected.
Interfaces	LLM.complete(prompt, model, temperature), Enrichment.lookup(handle domain), Social.publish(job), Social.metrics(platform_ref) — all as abstract classes with a fake implementation for tests.
Done when	A seeded company appears in Metabase; a job moves pending → executed via the queue; CI runs migrations and tests on every push.

Component 1 — Intake and ICP model weeks 1–2

Inputs	OAuth for LinkedIn and X; website URL, docs URL, optional CSV/API of product data; three best customers (company + person + why they bought); win/loss notes (optional).
Processing	Crawl site and docs → product summary (LLM). Enrich the three customers → firmographic pattern → ICP description + personas (LLM, editable). Embed ICP description and each customer.
Outputs	company, founder, icp rows; an editable ICP page in the app.
Done when	A new founder reaches a complete, edited ICP in under 20 minutes; the ICP embedding scores the three source customers above 0.8 similarity.

Component 2 — Scorecard (Judgment v0) week 2

Inputs	ICP, founder account stats (followers, 30-day impressions, post cadence), site crawl (pages, indexed count from Search Console if connected), launch history (Product Hunt scrape), signup flow (headless browser walkthrough).
Processing	Rule-based scoring 0–10 per channel: founder distribution, ICP attention, search, launches, onboarding. Each rule is a named function with a rationale string. One fix per channel chosen from a fix library by rule. LLM writes the narrative only.
Outputs	plan row with scorecard and a 90-day sequence (which channels in which order). Rendered scorecard page; exportable PDF (this is the Diagnostic deliverable).
Done when	Two humans reviewing the same company agree with the scorecard on 4 of 5 channels; the Simplicity diagnostic is produced by this component, not by hand.

Component 3 — Voice engine and approval feed weeks 3–5

Inputs	Weekly 45-minute interview (audio upload or call recording → transcript); founder's past posts (pulled once at intake); voice profile.
Processing	1) Extract 15–25 claims, stories and numbers from the transcript. 2) Map each to one of 12 formats (contrarian take, customer story, build log, teardown, number-with-lesson, before/after, framework, question, prediction, mistake, list, quote-with-take). 3) Draft per platform with the voice profile as system context. 4) Voice-match check: embedding similarity to the founder's own samples plus a banned-phrase and slop-pattern filter; drafts below threshold are regenerated, never shown. 5) Queue as Jobs with a scheduled slot.
Approval feed	Mobile-first list; tap to approve, swipe to reject, inline edit. Every edit is stored as a diff on the Job and appended to the voice profile samples. Approval, edit and rejection rates per format are tracked.
Publishing	Worker publishes at scheduled time via official APIs; stores platform_ref; retries with backoff; failure surfaces in the feed.
Done when	Three founders run two consecutive weeks with approval rate above 70% and fewer than 20% of posts edited; zero posts published without approval.

Component 4 — Attention map weeks 4–6

Inputs	ICP; the three customers' public follows and engagement; enrichment API; platform search.
Processing	Seed from the customers' networks and ICP keywords → expand via who-they-follow and who-engages → enrich each account → score ICP fit (embedding of headline + company vs ICP, weighted by role). Cluster into: buyers, influencers the buyers follow, peers, communities/newsletters.
Outputs	account rows; a daily list of 10 threads/posts by high-fit accounts with a drafted reply each (voice engine), delivered through the approval feed.
Done when	300+ scored accounts per company; founder rates 8 of 10 daily targets as "right people"; replies approved above 60%.

Component 5 — Metrics and Friday report weeks 5–7

Inputs	Per-job analytics from LinkedIn and X; follower lists; Search Console; signup webhook or embeddable field.
Impressions inside ICP	Definition: impressions × share of engaged/viewing accounts that match the ICP above a threshold, computed per job from available engagement data (reactors, commenters, profile-visit sources where the API exposes them) and extrapolated where it does not; the method is versioned and shown in the report. New followers are matched to the ICP the same way.
Outputs	Daily metric rows; Friday email and in-app report (this week vs last, per channel, top 3 posts, next week's plan); a public share card (PNG + page); the signup_source widget (one line of JS, or a Zapier/webhook receiver).
Done when	Report goes out automatically Friday 9am local for every active company; number is reproducible from stored rows; two founders share the card

unprompted.

Component 6 — Sequencer (Judgment v1) weeks 7–9

Inputs	Last week's outcome row; the 90-day sequence from the scorecard; approval and edit rates; content calendar constraints.
Processing	Rules, versioned in the repo: e.g. <i>if impressions_icp grew < 10% two weeks running and approval > 70%, shift 30% of slots to the top-performing format; if search score < 4 and product data exists, schedule the first page batch; if a launch is within 6 weeks, allocate launch tasks</i> . Emits the week's Jobs. Human override available in an internal admin; every override logged as a labeled example.
Done when	Monday plans for all companies are generated without a human for two consecutive weeks; overrides fall below 20% of decisions.

Component 7 — Page factory weeks 8–12

Templates	(a) <i>[product]</i> for <i>[segment]</i> ; (b) <i>[competitor]</i> alternative; (c) <i>[use case]</i> with <i>[data point]</i> from the company's own data. Each template: keyword rule, data requirements, section layout, internal-link rule, schema.org markup.
Processing	Keyword clustering from ICP and Search Console → candidate pages ranked by opportunity → data pulled from product sources → page generated with a real data point on every page (no data, no page) → founder approves the batch → published on {company}.pages.ourdomain or pushed via CMS adapter (Webflow, Framer, Next.js export) → sitemap and indexing requests → indexing tracked.
Done when	One company has 50 pages published and 30 indexed within 30 days; Search Console impressions from those pages appear in the Friday report.

Component 8 — Launch kit weeks 13–18

Processing	Launch object with a date and type (Product Hunt, feature drop, joint) → six-week calendar generated as Jobs (teasers, hunter outreach, asset creation, day-of sequence, follow-ups) → checklists in the app → assets drafted by the voice engine → metrics rolled into the report with a launch view.
Done when	Two companies complete a launch with every step tracked as a Job; launch-week signups attributed via <code>signup_source</code> .

Component 9 — Relationship engine v1 weeks 19–26 · v2 weeks 36–42

v1 (brokered launches)	Match companies on the product by adjacent ICP and non-competing product; propose a joint launch to both founders in-app; on mutual accept, generate the shared calendar and assets; track results for both.
v2 (outreach)	From the attention map's influencer and peer clusters, propose 5 relationships a month; draft warm outreach in the founder's voice (comment-first, then DM/email) via the feed; track replies and follow-ups as Jobs.

Done when v1: two joint launches completed between companies on the product. v2: reply rate on outreach above 25% across 50 sends.

Component 10 — Site and onboarding engine weeks 27–34

Processing Headless-browser audit against 40 checks (hero clarity, proof, CTA, signup steps, first-run, activation email) → scored → rewrites generated (copy, section order, email sequence) → founder approves → shipped via CMS adapter or served as a hosted variant with a redirect → signup-to-activation measured before/after.

Done when Three companies ship a rewritten flow; median activation rate improves versus their own baseline.

Component 11 — Tool factory weeks 32–38

Processing Sequencer proposes 3 tool ideas per quarter from product data + ICP questions (calculator, scorecard, generator, lookup) → founder picks one → product generates a single-page app from a template set (inputs, logic, result, lead capture) → hosted on the pages subdomain → distribution Jobs (posts, replies, page links).

Done when Five tools live across companies; each captures 50+ emails in its first month.

Component 12 — Learned judgment (v2) weeks 40–52

Prerequisite 100 companies × ≥ 8 weeks of outcome rows, with plan features and result deltas.

Processing Retrieval over similar companies' outcomes (pgvector on ICP + stage + channel state) feeding an LLM planner, then a small learned ranker over candidate sequences; blind A/B against the rules engine on new companies.

Done when v2 beats the rules engine on 4-week impressions-inside-ICP growth with $p < 0.05$; then it becomes the default with rules as fallback.

7. Quality gates and metrics

Metric	Definition	Gate
Approval rate	$\text{approved} \div (\text{approved} + \text{rejected})$ per founder per week, posts and replies separately	> 70% for 4 weeks before humans leave the loop for that founder
Edit rate	$\text{edited} \div \text{approved}$	< 20%
Impressions inside ICP	Weekly, method-versioned (Component 5)	3× baseline within 4 weeks on at least half of design partners by day 90
Followers matching ICP	New followers with ICP match > threshold	Reported; no gate in the wedge
Pages indexed	Search Console coverage	60% of published pages within 30 days
Founder hours per week	Interview + approvals, measured in-app	< 1.5 hours
Publish failures	Jobs in failed after retries	< 1% of executed
Human override rate	Sequencer decisions changed by staff	< 20% by week 9

Kill signal, restated for engineering. If by week 16 fewer than half of active companies show a 2× impressions-inside-ICP improvement with approval above 60%, stop adding components and narrow: one vertical, one channel, until the number moves.

8. Environment, security and cost

- **Environments:** local (Docker Compose: Postgres, Redis, API, worker, web), staging, production. Migrations run in CI; no manual schema changes.
- **Secrets:** OAuth tokens encrypted at rest (KMS-managed key); never logged; rotated on disconnect. Founders can revoke from the app; revocation cancels all pending Jobs.
- **Platform compliance:** official APIs only; publish only founder-approved content; rate limits respected per platform; no scraping of logged-in surfaces. The attention map uses public data and the enrichment vendor's licensed data.
- **Data ownership:** every company can export pages (HTML), posts, accounts and metrics as CSV/JSON in one click; deletion removes all rows and files within 24 hours.
- **Model governance:** prompts and model versions pinned per component; a change to a voice prompt is a versioned deploy with a canary on two founders.

- **Cost per company per month (estimate at 100 companies):** inference \$40–90, enrichment \$20–40, hosting and search APIs \$10–25, platform fees \$0. Target < \$150 all-in at the Team tier, consistent with the 80%+ margin target once human review is spot-check only.

9. First two weeks, day by day

Day	Work	Output
Mon 9/28	Repo, Compose, Postgres + Redis, schema package with 7 core tables, migrations, CI	Green CI, Metabase shows empty tables
Tue 9/29	FastAPI auth, /companies, /jobs; RQ worker; LLM/Enrichment/Social interfaces with fakes	Heartbeat job runs end to end
Wed 9/30	Next.js shell, login, company create; LinkedIn and X OAuth flows	A founder connects both accounts
Thu 10/1	Intake: site crawl → product summary; three-customer form → enrichment → ICP draft	Editable ICP page
Fri 10/2	ICP embeddings; pull founder's past posts; store voice samples	Component 1 done; Simplicity entered as company #1
Mon 10/5	Scorecard rules: founder distribution, ICP attention (from account stats + past posts)	Two channels scored
Tue 10/6	Scorecard rules: search (Search Console + crawl), launches, onboarding (headless walkthrough)	Five channels scored
Wed 10/7	Fix library; 90-day sequence generator; LLM narrative; scorecard page	Scorecard renders for Simplicity
Thu 10/8	PDF export of the scorecard; review with a second reader; adjust rules	Diagnostic PDF generated by the product
Fri 10/9	Send the Simplicity diagnostic to Juraj; start Component 3 (transcript → claims)	Component 2 done; first design-partner conversation

Sources for market and role data referenced by this plan are in the companion product plan. This document is the engineering sequence and is expected to change as the first components ship; every change to a definition of done is recorded in the repo's `PLAN.md`.